

# PySpark for BigMart Sales

In [1]:

```
from pyspark.sql import SparkSession
spark = SparkSession.builder.appName('Analysis-BigMart').getOrCreate()
data = spark.read.csv('train.csv', header = True, inferSchema = True)
data.printSchema()
```

```
root
 |-- Item_Identifier: string (nullable = true)
 |-- Item_Weight: double (nullable = true)
 |-- Item_Fat_Content: string (nullable = true)
 |-- Item_Visibility: double (nullable = true)
 |-- Item_Type: string (nullable = true)
 |-- Item_MRP: double (nullable = true)
 |-- Outlet_Identifier: string (nullable = true)
 |-- Outlet_Establishment_Year: integer (nullable = true)
 |-- Outlet_Size: string (nullable = true)
 |-- Outlet_Location_Type: string (nullable = true)
 |-- Outlet_Type: string (nullable = true)
 |-- Item_Outlet_Sales: double (nullable = true)
```

In [2]:

```
data.show(5, truncate=False)
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|Item_Identifier|Item_Weight|Item_Fat_Content|Item_Visibility|Item_Type|Item_MRP|Outlet_Identifier|Outlet_Establishment_Year|Outlet_Size|Outlet_Location_Type|Outlet_Type|Item_Outlet_Sales|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|FDA15|9.3|Low Fat|0.016047301|Dairy|249.8092|OUT049|1999|M|
edium|Tier 1|Supermarket Type1|3735.138|48.2692|OUT018|2009|M|
|DRC01|5.92|Regular|0.019278216|Soft Drinks|141.618|OUT049|1999|M|
edium|Tier 3|Supermarket Type2|443.4228|2097.27|OUT010|1998|n|
|FDN15|17.5|Low Fat|0.016760075|Meat|53.8614|OUT013|1987|H|
edium|Tier 1|Supermarket Type1|2097.27|732.38|
|FDX07|19.2|Regular|0.0|Fruits and Vegetables|1994.7052|
ull|Tier 3|Grocery Store|732.38|
|NCD19|8.93|Low Fat|0.0|Household|
igh|Tier 3|Supermarket Type1|994.7052|
only showing top 5 rows
```

In [3]:

```
data.describe().show()
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|summary|Item_Identifier|Item_Weight|Item_Fat_Content|Item_Visibility|Item_Type|Item_MRP|Outlet_Identifier|Outlet_Establishment_Year|Outlet_Size|Outlet_Location_Type|Outlet_Type|Item_Outlet_Sales|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|count|8523|6113|8523|8523|8523|8523|8523|8523|
|mean|null|12.857645184136183|null|0.06613202877895127|null|140.9927819781768|null|
1997.8318667135984|null|4.643456499186414|null|0.05159782232113514|null|62.27506651219047|null|
|stddev|8.371760408092667|null|4.555|null|1706.4996157338403|
|min|DRA12|4.555|LF|0.0|Baking Goods|31.29|OUT010|
1985|High|Tier 1|Grocery Store|33.29|
|max|NCZ54|21.35|reg|0.328390948|Starchy Foods|266.8884|OUT049|
2009|Small|Tier 3|Supermarket Type3|13086.9648|
```

In [4]:

```
import pyspark.sql.functions as f

# null values in each column
data_agg = data.agg(*[f.count(f.when(f.isnull(c), c)).alias(c) for c in data.columns])
data_agg.show()
```

Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Establishment_Year	Outlet_Size	Outlet_Location_Type	Outlet_Type	Item_Outlet_Sales
0	1463	0	0	0	0	0	0	2410			

## Fill NA with mean and mode

In [5]:

```
data.agg({'Item_Weight': 'mean'}).show()
```

avg(Item_Weight)
12.857645184136183

In [6]:

```
data = data.fillna( { 'Item_Weight':12.857645184136183 } )
```

In [7]:

```
data.groupBy('Outlet_Size').count().show()
```

Outlet_Size	count
High	932
null	2410
Medium	2793
Small	2388

In [8]:

```
data = data.fillna( { 'Outlet_Size':'Medium' } )
```

In [9]:

```
import pyspark.sql.functions as f

# null values in each column
data_agg = data.agg(*[f.count(f.when(f.isnull(c), c)).alias(c) for c in data.columns])
data_agg.show()
```

Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Establishment_Year	Outlet_Size	Outlet_Location_Type	Outlet_Type	Item_Outlet_Sales
0	0	0	0	0	0	0	0	0			

## Label Encoding

In [10]:

```

from pyspark.ml.feature import StringIndexer

# create object of StringIndexer class and specify input and output column
SI_Fat = StringIndexer(inputCol='Item_Fat_Content',outputCol='Item_Fat_Content_Index')
SI_Type = StringIndexer(inputCol='Item_Type',outputCol='Item_Type_Index')
SI_Size = StringIndexer(inputCol='Outlet_Size',outputCol='Outlet_Size_Index')
SI_Location = StringIndexer(inputCol='Outlet_Location_Type',outputCol='Outlet_Location_Type_Index')
SI_Out_Type = StringIndexer(inputCol='Outlet_Type',outputCol='Outlet_Type_Index')

# transform the data
data = SI_Fat.fit(data).transform(data)
data = SI_Type.fit(data).transform(data)
data = SI_Size.fit(data).transform(data)
data = SI_Location.fit(data).transform(data)
data = SI_Out_Type.fit(data).transform(data)

# view the transformed data
data.first()

```

```

Row(Item_Identifier='FDA15', Item_Weight=9.3, Item_Fat_Content='Low Fat', Item_Visibility=0.016047301, Item_Type='Dairy', Item_MRP=249.809
2, Outlet_Identifier='OUT049', Outlet_Establishment_Year=1999, Outlet_Size='Medium', Outlet_Location_Type='Tier 1', Outlet_Type='Supermark
et Type1', Item_Outlet_Sales=3735.138, Item_Fat_Content_Index=0.0, Item_Type_Index=4.0, Outlet_Size_Index=0.0, Outlet_Location_Type_Index=
2.0, Outlet_Type_Index=0.0)

```

## One Hot Encoding

In [11]:

```

from pyspark.ml.feature import OneHotEncoder
OHE = OneHotEncoder(inputCols=['Item_Fat_Content_Index', 'Item_Type_Index', 'Outlet_Size_Index', 'Outlet_
Location_Type_Index', 'Outlet_Type_Index'],outputCols=['Item_Fat_Content_OHE', 'Item_Type_OHE', 'Outlet_S
ize_OHE', 'Outlet_Location_Type_OHE', 'Outlet_Type_OHE'])

# transform the data
data = OHE.fit(data).transform(data)

```

In [12]:

```
data.take(1)
```

```

[Row(Item_Identifier='FDA15', Item_Weight=9.3, Item_Fat_Content='Low Fat', Item_Visibility=0.016047301, Item_Type='Dairy', Item_MRP=249.80
92, Outlet_Identifier='OUT049', Outlet_Establishment_Year=1999, Outlet_Size='Medium', Outlet_Location_Type='Tier 1', Outlet_Type='Supermar
ket Type1', Item_Outlet_Sales=3735.138, Item_Fat_Content_Index=0.0, Item_Type_Index=4.0, Outlet_Size_Index=0.0, Outlet_Location_Type_Index
=2.0, Outlet_Type_Index=0.0, Outlet_Type_OHE=SparseVector(3, {0: 1.0}), Item_Fat_Content_OHE=SparseVector(4, {0: 1.0}), Outlet_Size_OHE=Sp
arseVector(2, {0: 1.0}), Outlet_Location_Type_OHE=SparseVector(2, {}), Item_Type_OHE=SparseVector(15, {4: 1.0}))]

```

## Vectorizing into single Feature

In [13]:

```

from pyspark.ml.feature import VectorAssembler

inputcol = ['Item_Fat_Content_Index', 'Item_Type_Index', 'Outlet_Size_Index', 'Outlet_Location_Type_Inde
x', 'Outlet_Type_Index', 'Item_Fat_Content_OHE', 'Item_Type_OHE', 'Outlet_Size_OHE', 'Outlet_Location_Type
_OHE', 'Outlet_Type_OHE', 'Item_Weight', 'Item_Visibility', 'Item_MRP', 'Outlet_Establishment_Year']
# specify the input and output columns of the vector assembler
assembler = VectorAssembler(inputCols=inputcol, outputCol='features')

# fill the null values
data = data.fillna(0)

# transform the data
final_data = assembler.transform(data)

```

In [14]:

```
final_data.take(1)
```

```
[Row(Item_Identifier='FDA15', Item_Weight=9.3, Item_Fat_Content='Low Fat', Item_Visibility=0.016047301, Item_Type='Dairy', Item_MRP=249.8092, Outlet_Identifier='OUT049', Outlet_Establishment_Year=1999, Outlet_Size='Medium', Outlet_Location_Type='Tier 1', Outlet_Type='Supermarket Type1', Item_Outlet_Sales=3735.138, Item_Fat_Content_Index=0.0, Item_Type_Index=4.0, Outlet_Size_Index=0.0, Outlet_Location_Type_Index=2.0, Outlet_Type_Index=0.0, Outlet_Type_OHE=SparseVector(3, {0: 1.0}), Item_Fat_Content_OHE=SparseVector(4, {0: 1.0}), Outlet_Size_OHE=SparseVector(2, {0: 1.0}), Outlet_Location_Type_OHE=SparseVector(2, {}), Item_Type_OHE=SparseVector(15, {4: 1.0}), features=SparseVector(35, {1: 4.0, 3: 2.0, 5: 1.0, 13: 1.0, 24: 1.0, 28: 1.0, 31: 9.3, 32: 0.016, 33: 249.8092, 34: 1999.0}))]
```

In [15]:

```
df_train = final_data.select(['features', 'Item_Outlet_Sales'])
```

In [16]:

```
df_train.take(1)
```

```
[Row(features=SparseVector(35, {1: 4.0, 3: 2.0, 5: 1.0, 13: 1.0, 24: 1.0, 28: 1.0, 31: 9.3, 32: 0.016, 33: 249.8092, 34: 1999.0}), Item_Outlet_Sales=3735.138)]
```

In [17]:

```
df_train.show(5)
```

```
+-----+-----+
|          features|Item_Outlet_Sales|
+-----+-----+
|(35,[1,3,5,13,24,...]|      3735.138|
|(35,[0,1,4,6,17,2...]|      443.4228|
|(35,[1,3,5,18,24,...]|       2097.27|
|(35,[0,4,6,9,24,2...]|        732.38|
|(35,[1,2,5,11,26,...]|      994.7052|
+-----+-----+
only showing top 5 rows
```

## Split train and test

In [18]:

```
train_df, test_df = df_train.randomSplit([0.8, 0.2])
```

## Applying various models in Pyspark Api

```

In [30]:
# Linear Regression Model
from pyspark.ml.regression import LinearRegression
from pyspark.ml.evaluation import RegressionEvaluator

lr = LinearRegression(featuresCol = 'features', labelCol='Item_Outlet_Sales', maxIter=150, regParam=0.0,
elasticNetParam=0.0)
lr_model = lr.fit(train_df)

trainingSummary = lr_model.summary
print("RMSE: %f" % trainingSummary.rootMeanSquaredError)
print("R2: %f" % trainingSummary.r2)

# Predict and Evaluate
lr_predictions = lr_model.transform(test_df)
lr_predictions.select("prediction", "Item_Outlet_Sales", "features").show(5)

lr_evaluator = RegressionEvaluator(predictionCol="prediction", labelCol="Item_Outlet_Sales", metricName="rmse")
print("Root Mean Square Error (RMSE) on test data = %g" % lr_evaluator.evaluate(lr_predictions))

lr_evaluator_r2 = RegressionEvaluator(predictionCol="prediction", labelCol="Item_Outlet_Sales", metricName="r2")
print("R Squared (R2) on test data = %g" % lr_evaluator_r2.evaluate(lr_predictions))

RMSE: 1133.476559
R2: 0.554448
+-----+-----+-----+
| prediction|Item_Outlet_Sales| features|
+-----+-----+-----+
| 729.2843759230018| 331.5684|(35,[0,1,2,3,4,6,...|
| 384.5975328735367| 425.4462|(35,[0,1,2,3,4,6,...|
| 1136.7232229275978| 945.436|(35,[0,1,2,3,4,6,...|
|-153.25180773547618| 317.5866|(35,[0,1,2,3,4,6,...|
| -30.22022594201553| 113.8518|(35,[0,1,2,3,4,6,...|
+-----+-----+-----+
only showing top 5 rows

Root Mean Square Error (RMSE) on test data = 1106.76
R Squared (R2) on test data = 0.593892

```

```

In [23]:
# Lasso Regression Model
lar = LinearRegression(featuresCol = 'features', labelCol='Item_Outlet_Sales', maxIter=100, regParam=0.3,
elasticNetParam=1)
lar_model = lar.fit(train_df)

trainingSummary = lar_model.summary
print("RMSE: %f" % trainingSummary.rootMeanSquaredError)
print("R2: %f" % trainingSummary.r2)

# Predict and Evaluate
lar_predictions = lar_model.transform(test_df)
lar_predictions.select("prediction", "Item_Outlet_Sales", "features").show(5)

lar_evaluator = RegressionEvaluator(predictionCol="prediction", labelCol="Item_Outlet_Sales", metricName=
"rmse")
print("Root Mean Square Error (RMSE) on test data = %g" % lar_evaluator.evaluate(lar_predictions))

lar_evaluator_r2 = RegressionEvaluator(predictionCol="prediction", labelCol="Item_Outlet_Sales", metricName=
"r2")
print("R Squared (R2) on test data = %g" % lar_evaluator_r2.evaluate(lar_predictions))

RMSE: 1133.589738
R2: 0.554359
+-----+-----+-----+
| prediction|Item_Outlet_Sales| features|
+-----+-----+-----+
| 723.1792481893208| 331.5684|(35,[0,1,2,3,4,6,...|
| 378.63698723577545| 425.4462|(35,[0,1,2,3,4,6,...|
| 1130.646491530053| 945.436|(35,[0,1,2,3,4,6,...|
|-158.85101700072846| 317.5866|(35,[0,1,2,3,4,6,...|
| -35.82258349095355| 113.8518|(35,[0,1,2,3,4,6,...|
+-----+-----+-----+
only showing top 5 rows

Root Mean Square Error (RMSE) on test data = 1106.86
R Squared (R2) on test data = 0.593819

```

In [26]:

```
# Decision Tree Regression Model
from pyspark.ml.regression import DecisionTreeRegressor

dt = DecisionTreeRegressor(featuresCol = 'features', labelCol='Item_Outlet_Sales', maxDepth=15, minInstancesPerNode=100)
dt_model = dt.fit(train_df)

# Predict and Evaluate
dt_predictions = dt_model.transform(test_df)
dt_predictions.select("prediction", "Item_Outlet_Sales", "features").show(5)

df_evaluator = RegressionEvaluator(predictionCol="prediction", labelCol="Item_Outlet_Sales", metricName="rmse")
print("Root Mean Square Error (RMSE) on test data = %g" % df_evaluator.evaluate(dt_predictions))

df_evaluator_r2 = RegressionEvaluator(predictionCol="prediction", labelCol="Item_Outlet_Sales", metricName="r2")
print("R Squared (R2) on test data = %g" % df_evaluator_r2.evaluate(dt_predictions))
```

prediction	Item_Outlet_Sales	features
394.03078252427184	331.5684	(35, [0, 1, 2, 3, 4, 6, ...]
283.3360183206108	425.4462	(35, [0, 1, 2, 3, 4, 6, ...]
500.69881896551715	945.436	(35, [0, 1, 2, 3, 4, 6, ...]
223.94658571428573	317.5866	(35, [0, 1, 2, 3, 4, 6, ...]
283.3360183206108	113.8518	(35, [0, 1, 2, 3, 4, 6, ...]

only showing top 5 rows

Root Mean Square Error (RMSE) on test data = 1057.63  
R Squared (R2) on test data = 0.629148

In [27]:

```
## GBT Regression Model
from pyspark.ml.regression import GBTRegressor

gbt = GBTRegressor(featuresCol = 'features', labelCol='Item_Outlet_Sales', maxIter=10)
# gbt = GBTRegressor(featuresCol = 'features', labelCol='Item_Outlet_Sales', maxIter=10, maxDepth=15, minInstancesPerNode=100)
gbt_model = gbt.fit(train_df)

# Predict and Evaluate
gbt_predictions = gbt_model.transform(test_df)
gbt_predictions.select("prediction", "Item_Outlet_Sales", "features").show(5)

gbt_evaluator = RegressionEvaluator(predictionCol="prediction", labelCol="Item_Outlet_Sales", metricName="rmse")
print("Root Mean Square Error (RMSE) on test data = %g" % gbt_evaluator.evaluate(gbt_predictions))

gbt_evaluator_r2 = RegressionEvaluator(predictionCol="prediction", labelCol="Item_Outlet_Sales", metricName="r2")
print("R Squared (R2) on test data = %g" % gbt_evaluator_r2.evaluate(gbt_predictions))
```

prediction	Item_Outlet_Sales	features
477.42167799114003	331.5684	(35, [0, 1, 2, 3, 4, 6, ...]
342.25790791842667	425.4462	(35, [0, 1, 2, 3, 4, 6, ...]
541.6654409408936	945.436	(35, [0, 1, 2, 3, 4, 6, ...]
245.1703909636694	317.5866	(35, [0, 1, 2, 3, 4, 6, ...]
311.27774749352017	113.8518	(35, [0, 1, 2, 3, 4, 6, ...]

only showing top 5 rows

Root Mean Square Error (RMSE) on test data = 1047.03  
R Squared (R2) on test data = 0.636544

```

In [29]:
# RandomForest Regression Model
from pyspark.ml.tuning import CrossValidator, ParamGridBuilder
from pyspark.ml.regression import RandomForestRegressor

rf = RandomForestRegressor(featuresCol = 'features', labelCol='Item_Outlet_Sales', minInstancesPerNode=15
0)

rf_evaluator = RegressionEvaluator(predictionCol="prediction", labelCol="Item_Outlet_Sales",metricName="r
mse")

paramGrid = ParamGridBuilder() \
    .addGrid(rf.maxDepth, [5, 6, 10]) \
    .addGrid(rf.numTrees, [200, 400]) \
    .build()

cv = CrossValidator(estimator=rf, estimatorParamMaps=paramGrid, evaluator=rf_evaluator, numFolds=5, paral
lelism=5)

cv_model = cv.fit(train_df)

rf_predictions = cv_model.transform(test_df)
rf_predictions.select("prediction", "Item_Outlet_Sales", "features").show(5)

print("Root Mean Square Error (RMSE) on test data = %g" % gbt_evaluator.evaluate(gbt_predictions))

rf_evaluator_r2 = RegressionEvaluator(predictionCol="prediction", labelCol="Item_Outlet_Sales",metricName
="r2")
print("R Squared (R2) on test data = %g" % rf_evaluator_r2.evaluate(rf_predictions))

```

```

+-----+-----+-----+
|      prediction|Item_Outlet_Sales|      features|
+-----+-----+-----+
|441.65932039560516|      331.5684|(35,[0,1,2,3,4,6,...|
| 325.619648765012|      425.4462|(35,[0,1,2,3,4,6,...|
|473.97048510979783|      945.436|(35,[0,1,2,3,4,6,...|
|309.47752994468544|      317.5866|(35,[0,1,2,3,4,6,...|
|319.80938730928233|      113.8518|(35,[0,1,2,3,4,6,...|
+-----+-----+-----+
only showing top 5 rows

```

```


Root Mean Square Error (RMSE) on test data = 1047.03
R Squared (R2) on test data = 0.61196

```

## Results



```
In [33]:  
  
import numpy as np  
import matplotlib.pyplot as plt  
  
models = ['LinearRegression', 'LassoRegression', 'DecisionTreeRegressor',  
          'GBRegressor', 'RandomForestRegressor']  
  
data = [1106.76, 1106.86, 1057.63, 1047.03, 1047.03]  
  
plt.barh(models, data)  
  
for index, value in enumerate(data):  
    plt.text(value, index,  
             str(value))  
  
plt.show()
```



Here we can easily see that RFR and GBT outplayed all the other models as they have least RMSE scores on test\_dataset.

Improvements can be done as we can apply ANN and also can change the parameters in CrossValidations and other models' arguments as well.

We can even apply word embeddings as in case of neuralNets as well.